

ECCV 2012 Tutorial

Part 2

Explicit Embeddings I


Kernel Feature Maps

Andrea Vedaldi
University of Oxford


modelling of structure

massive datasets


location




direction



deformation



truncation



[Vedaldi Zisserman 09]



Challenge

1000 classes




1.5M images

>50k-dim. descriptors

Very efficient learning

bicycle?

\mathbf{x}




binary classifier

$$F(\mathbf{x}) = \langle \mathbf{w}, \mathbf{x} \rangle$$

Linear SVM


- ✓ fast
- ✗ restrictive



$$F(\mathbf{x}) = \langle \mathbf{w}, \mathbf{x} \rangle$$

Non-linear SVM

- ✗ much slower
- ✓ powerful



$$F(\mathbf{x}) = \sum_{i=1}^N \alpha_i K(\mathbf{x}, \mathbf{x}_i)$$

$$F(\mathbf{x}) = \sum_{i=1}^N \alpha_i K(\mathbf{x}, \mathbf{x}_i)$$

thousand bicycles



many more non-bicycle



$$F(\mathbf{x}) = \sum_{i=1}^N \alpha_i K(\mathbf{x}, \mathbf{x}_i)$$



feature map

$$K(\mathbf{x}, \mathbf{x}_i) = \langle \Psi(\mathbf{x}), \Psi(\mathbf{x}_i) \rangle$$

$$F(\mathbf{x}) = \langle \mathbf{w}, \Psi(\mathbf{x}) \rangle \qquad \mathbf{w} = \sum_{i=1}^N \alpha_i \Psi(\mathbf{x}_i)$$

$$F(\mathbf{x}) = \sum_{i=1}^N \alpha_i K(\mathbf{x}, \mathbf{x}_i)$$



approximated feature map

$$K(\mathbf{x}, \mathbf{x}_i) \approx \langle \hat{\Psi}(\mathbf{x}), \hat{\Psi}(\mathbf{x}_i) \rangle$$


$$F(\mathbf{x}) = \langle \mathbf{w}, \hat{\Psi}(\mathbf{x}) \rangle \qquad \mathbf{w} = \sum_{i=1}^N \alpha_i \hat{\Psi}(\mathbf{x}_i)$$

X^2 kernel

Excellent for bag-of-words, ...

$$k(x, x') = \frac{2xx'}{x + x'}$$

$$k(x, x')$$




Homogeneous kernel map


Closed form, simple, small

$$\Phi(x) = \sqrt{x} \begin{bmatrix} 0.8 \\ 0.6 \cos(0.6 \log x) \\ 0.6 \sin(0.6 \log x) \end{bmatrix}$$

$$\langle \Phi(x), \Phi(x') \rangle$$



$$\langle \Phi(x), \Phi(0.25) \rangle$$





[Vedaldi Zisserman 10,11]

Example: *homogeneous kernel map*

MATLAB code for Chi2 kernel

```
x = .01:.01:1 ;
for i = 1:100
    for j = 1:100
        K(i,j) = ...
            2*x(i)*x(j)/(x(i)+x(j));
    end
end
```




With the hom. kernel feature map

```
x = .01:.01:1 ;
psi = vl_homkernmap(x,1) ;
K = psi'*psi ;
```




VLFeat Toolbox
<http://www.vlfeat.org>



Caltech-101 category recognition



#1,500

training time
1 h  5 m
4× speedup

DaimlerChrysler pedestrian recognition



#20,000

1/2 h  14 s
100× speedup

Trecvid 2009 video indexing



#70,000

> 1 h  22.6 s
160× speedup

Finite-dimensional embeddings by projection

Exact feature

$$\Psi(\mathbf{x})$$


exact but inf. dim.

projection

Approximated feature

$$\hat{\Psi}(\mathbf{x})$$

approx. but compact:
small + dense
or sparse




Exact feature space

(reproducing kernel Hilbert space)

$$K(\mathbf{x}, \mathbf{x}') = \langle \Psi(\mathbf{x}), \Psi(\mathbf{x}') \rangle_{\mathcal{H}}$$

Data distribution

$$p(\mathbf{x})$$



PCA in feature space

Top D eigenfunctions of the kernel

$$\int_{\mathcal{X}} K(\mathbf{x}, \mathbf{z}) u_i(\mathbf{z}) p(\mathbf{z}) d\mathbf{z} = \kappa_i^2 u_i(\mathbf{x})$$

Coordinate functions

Projection on orthonormal PCA basis

$$\Phi_i(\mathbf{x}) = \kappa_i u_i(\mathbf{x})$$


Approximate feature

$$\hat{\Psi}(\mathbf{x}) \cong \Phi(\mathbf{x}) = \begin{bmatrix} \Phi_1(\mathbf{x}) \\ \vdots \\ \Phi_D(\mathbf{x}) \end{bmatrix}$$

Stationary kernel

Translation invariant


$$\forall \mathbf{t} \in \mathbb{R}^d : K(\mathbf{x} + \mathbf{t}, \mathbf{x}' + \mathbf{t}) = K(\mathbf{x}, \mathbf{x}')$$



Signature / profile


1D positive definite function

$$K(\mathbf{x}, \mathbf{x}') = \mathcal{K}(\mathbf{x}' - \mathbf{x})$$



Eigenfunctions = Sinusoids

$$\int_{\mathbb{R}^d} \mathcal{K}(\mathbf{x} - \mathbf{z}) e^{-i\langle \boldsymbol{\omega}, \mathbf{z} \rangle} d\mathbf{z} = \kappa_{\boldsymbol{\omega}}^2 e^{-i\langle \boldsymbol{\omega}, \mathbf{x} \rangle}$$



$$\Phi_{\boldsymbol{\omega}}(\mathbf{x}) = \kappa_{\boldsymbol{\omega}} e^{-i\langle \boldsymbol{\omega}, \mathbf{x} \rangle}$$

Fourier
transform



Additive kernel

Sum of 1D kernels

$$K(\mathbf{x}, \mathbf{x}') = \sum_{l=1}^D k(x_l, x'_l)$$



Hellinger

$$k(x, x') = \sqrt{xx'}$$





x^2

$$\frac{2xx'}{x + x'}$$




intersection

$$\min\{x, x'\}$$




The trick


$$k(x, x')$$



$$\frac{k(x, x')}{\sqrt{xx'}}$$



$$\log x$$



Homogeneous kernel

Multiplicative constant pops out

$$\forall c \geq 0 : k(cx, cx') = ck(x, x')$$

Signature / profile

Up to a factor and a logarithm

$$k(x, x') = \sqrt{xx'} \mathcal{K}(\log x - \log x')$$

$$\Phi_\omega(x) = \kappa_\omega \sqrt{x} e^{-\mathbf{i}\langle \omega, \log x \rangle}$$

Hellinger **X^2** **intersection**

$$k(x, x') = \sqrt{xx'}$$

$$\frac{2xx'}{x + x'}$$

$$\min\{x, x'\}$$

$$\mathcal{K}(\lambda) = 1$$

$$e^{-|\lambda|/2}$$

$$\operatorname{sech}(\lambda/2)$$

$$\kappa_{\omega}^2 = \delta(\omega)$$

$$\frac{2}{\pi(1 + 4\omega^2)}$$

$$\operatorname{sech}(\pi\omega)$$

$$\Phi_{\omega}(x) = \sqrt{x}$$

$$\sqrt{\frac{2x}{\pi(1 + 4\omega^2)}} e^{-\mathbf{i} \omega \log x}$$

$$\sqrt{x \operatorname{sech}(\pi\omega)} e^{-\mathbf{i} \omega \log x}$$


Exact feature space

(e.g. reproducing kernel Hilbert space)

$$K(\mathbf{x}, \mathbf{x}') = \langle \Psi(\mathbf{x}), \Psi(\mathbf{x}') \rangle_{\mathcal{H}}$$

Data distribution

$$\{\mathbf{x}_1, \dots, \mathbf{x}_n\}$$



PCA in feature space

Find the top D eigenfunctions of the kernel

$$\frac{1}{n} \sum_{j=1}^n K(\mathbf{x}_i, \mathbf{x}_j) u_i(\mathbf{x}_j) = \kappa_i^2 u_i(\mathbf{x}_i)$$


Coordinate functions

Projection on orthonormal PCA basis

$$\Phi_i(\mathbf{x}) = \kappa_i u_i(\mathbf{x})$$

$$= (n\kappa_i)^{-1} \sum_{j=1}^n K(\mathbf{x}, \mathbf{x}_j) u_i(\mathbf{x}_j)$$

encoding a new point \mathbf{x}
requires projecting it
on all the data sample



$$\Phi_i(x) = (n\kappa_i)^{-1} \sum_{j=1}^n k(x, x_j) u_i(x_j)$$


[Perronnin *et al.* 10]

- For any additive kernel
 - Empirical Nyström's approximation can be applied component-wise
 - The coordinate functions **can be tabulated**

- RBF / Gaussian / exponential kernel
Random Fourier features [Rahimi Recht 07]

$$Q(\mathbf{x}, \mathbf{x}') = \exp \left(-\frac{1}{2\sigma^2} \|\mathbf{x} - \mathbf{x}'\|^2 \right)$$

$$\Phi_{\text{RBF}}(\mathbf{x}) = D^{-\frac{1}{2}} \left[\cos\langle\omega_1, \mathbf{x}\rangle \quad \dots \quad \cos\langle\omega_D, \mathbf{x}\rangle \quad \sin\langle\omega_1, \mathbf{x}\rangle \quad \dots \right]^\top$$


 random Gaussian vectors

- **Generalized: RBF + Chi2 distance**

$$Q(\mathbf{x}, \mathbf{x}') = \exp \left(-\frac{1}{2\sigma^2} \sum_{i=1}^D \frac{2(x_i - x'_i)^2}{x_i + x'_i} \right)$$


$$\Phi_{\text{GRBF}}(\mathbf{x}) = D^{-\frac{1}{2}} \left[\cos\langle\omega_1, \Phi(\mathbf{x})\rangle \quad \dots \quad \cos\langle\omega_D, \Phi(\mathbf{x})\rangle \quad \sin\langle\omega_1, \Phi(\mathbf{x})\rangle \quad \dots \right]^\top$$

RBF kernel

```
[x1,x2] = meshgrid(range) ;
x = [x1(:) x2(:)]' ;
xp = [0.5;0.5] ;
for i=1:n*n
    sqd1 = (x(1,i) - xp(1))^2 ;

    sqd2 = (x(2,i) - xp(2))^2 ;


    K(i) = exp(-0.5*(sqd1 + sqd2)) ;
end
```



RBF w/Random Fourier Features


```
omega = randn(10,2) ;

psi = [cos(omega*x) ;
        sin(omega*x)] / sqrt(10) ;
psip = [cos(omega*xp) ;
         sin(omega*xp)] / sqrt(10) ;
K = psi'*psip ;
```




Chi2-RBF kernel

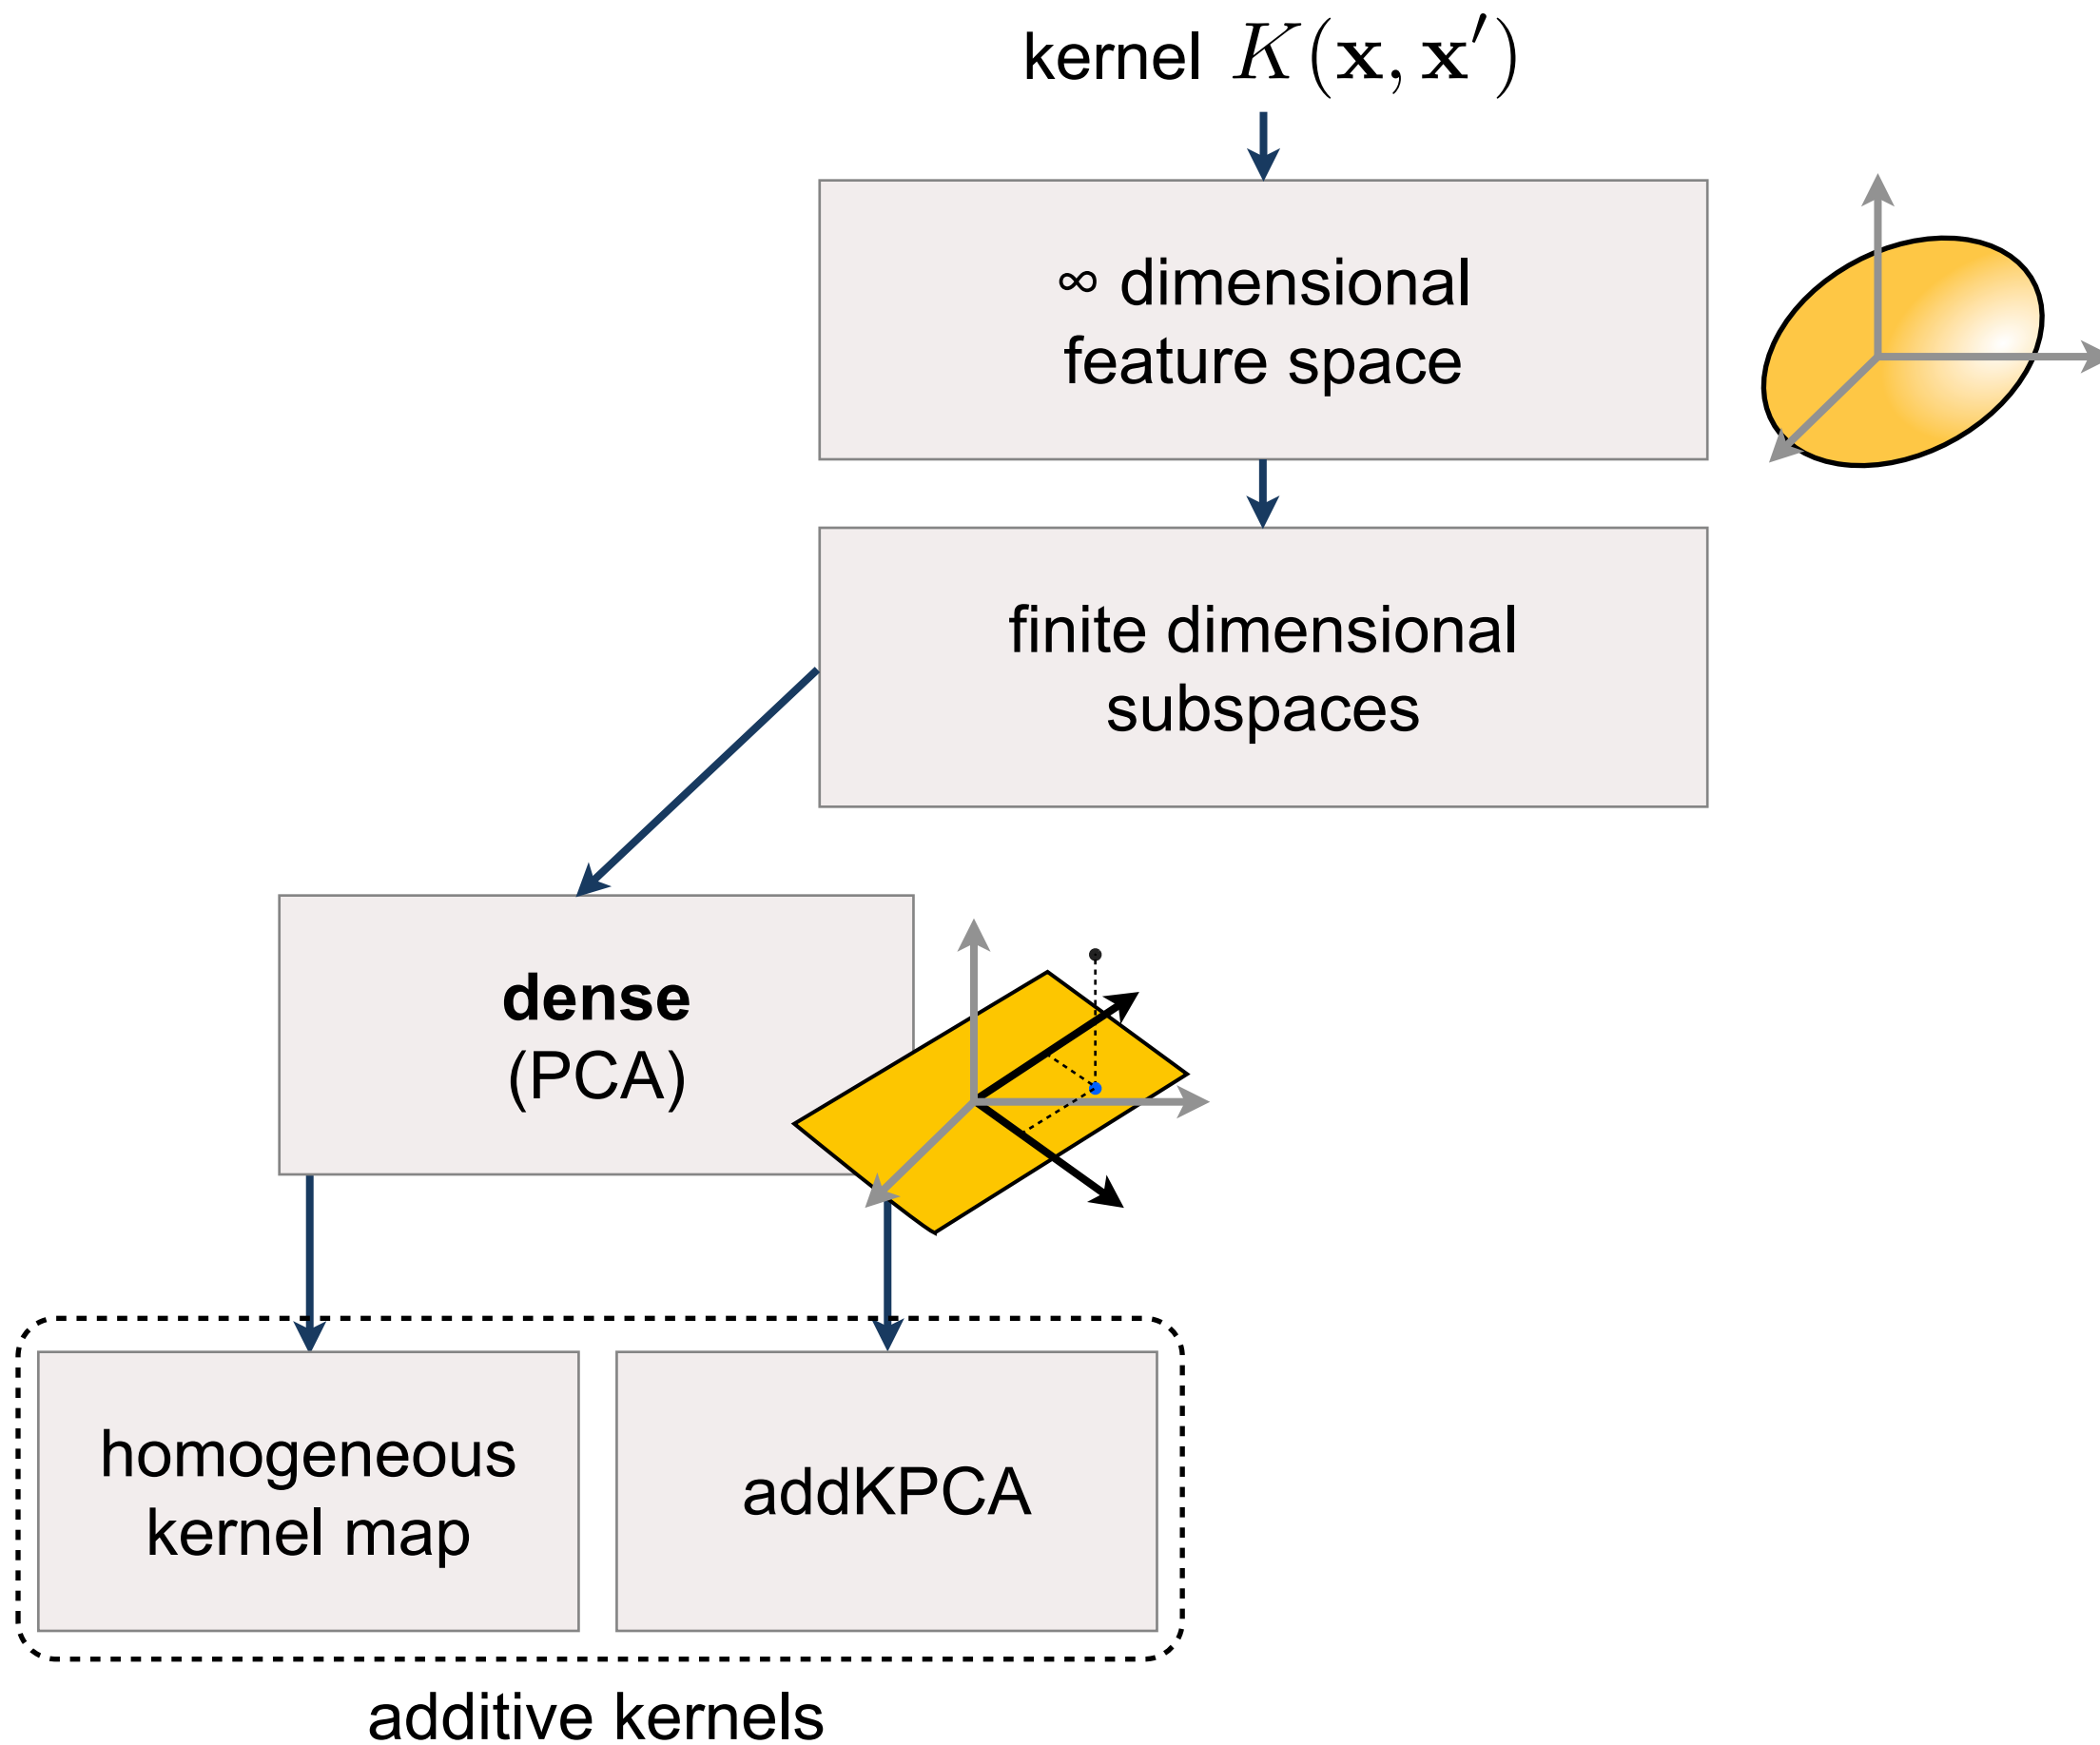
```
[x1,x2] = meshgrid(range) ;
x = [x1(:) x2(:)]' ;
xp = [0.5;0.5] ;
for i=1:n*n
    sqd1 = (x(1,i) - xp(1))^2 /
(x(1,i)+xp(1)) ;
    sqd2 = (x(2,i) - xp(2))^2 /
(x(2,i)+xp(2)) ;
    K(i) = exp(-0.5*(sqd1 + sqd2)) ;
end
```

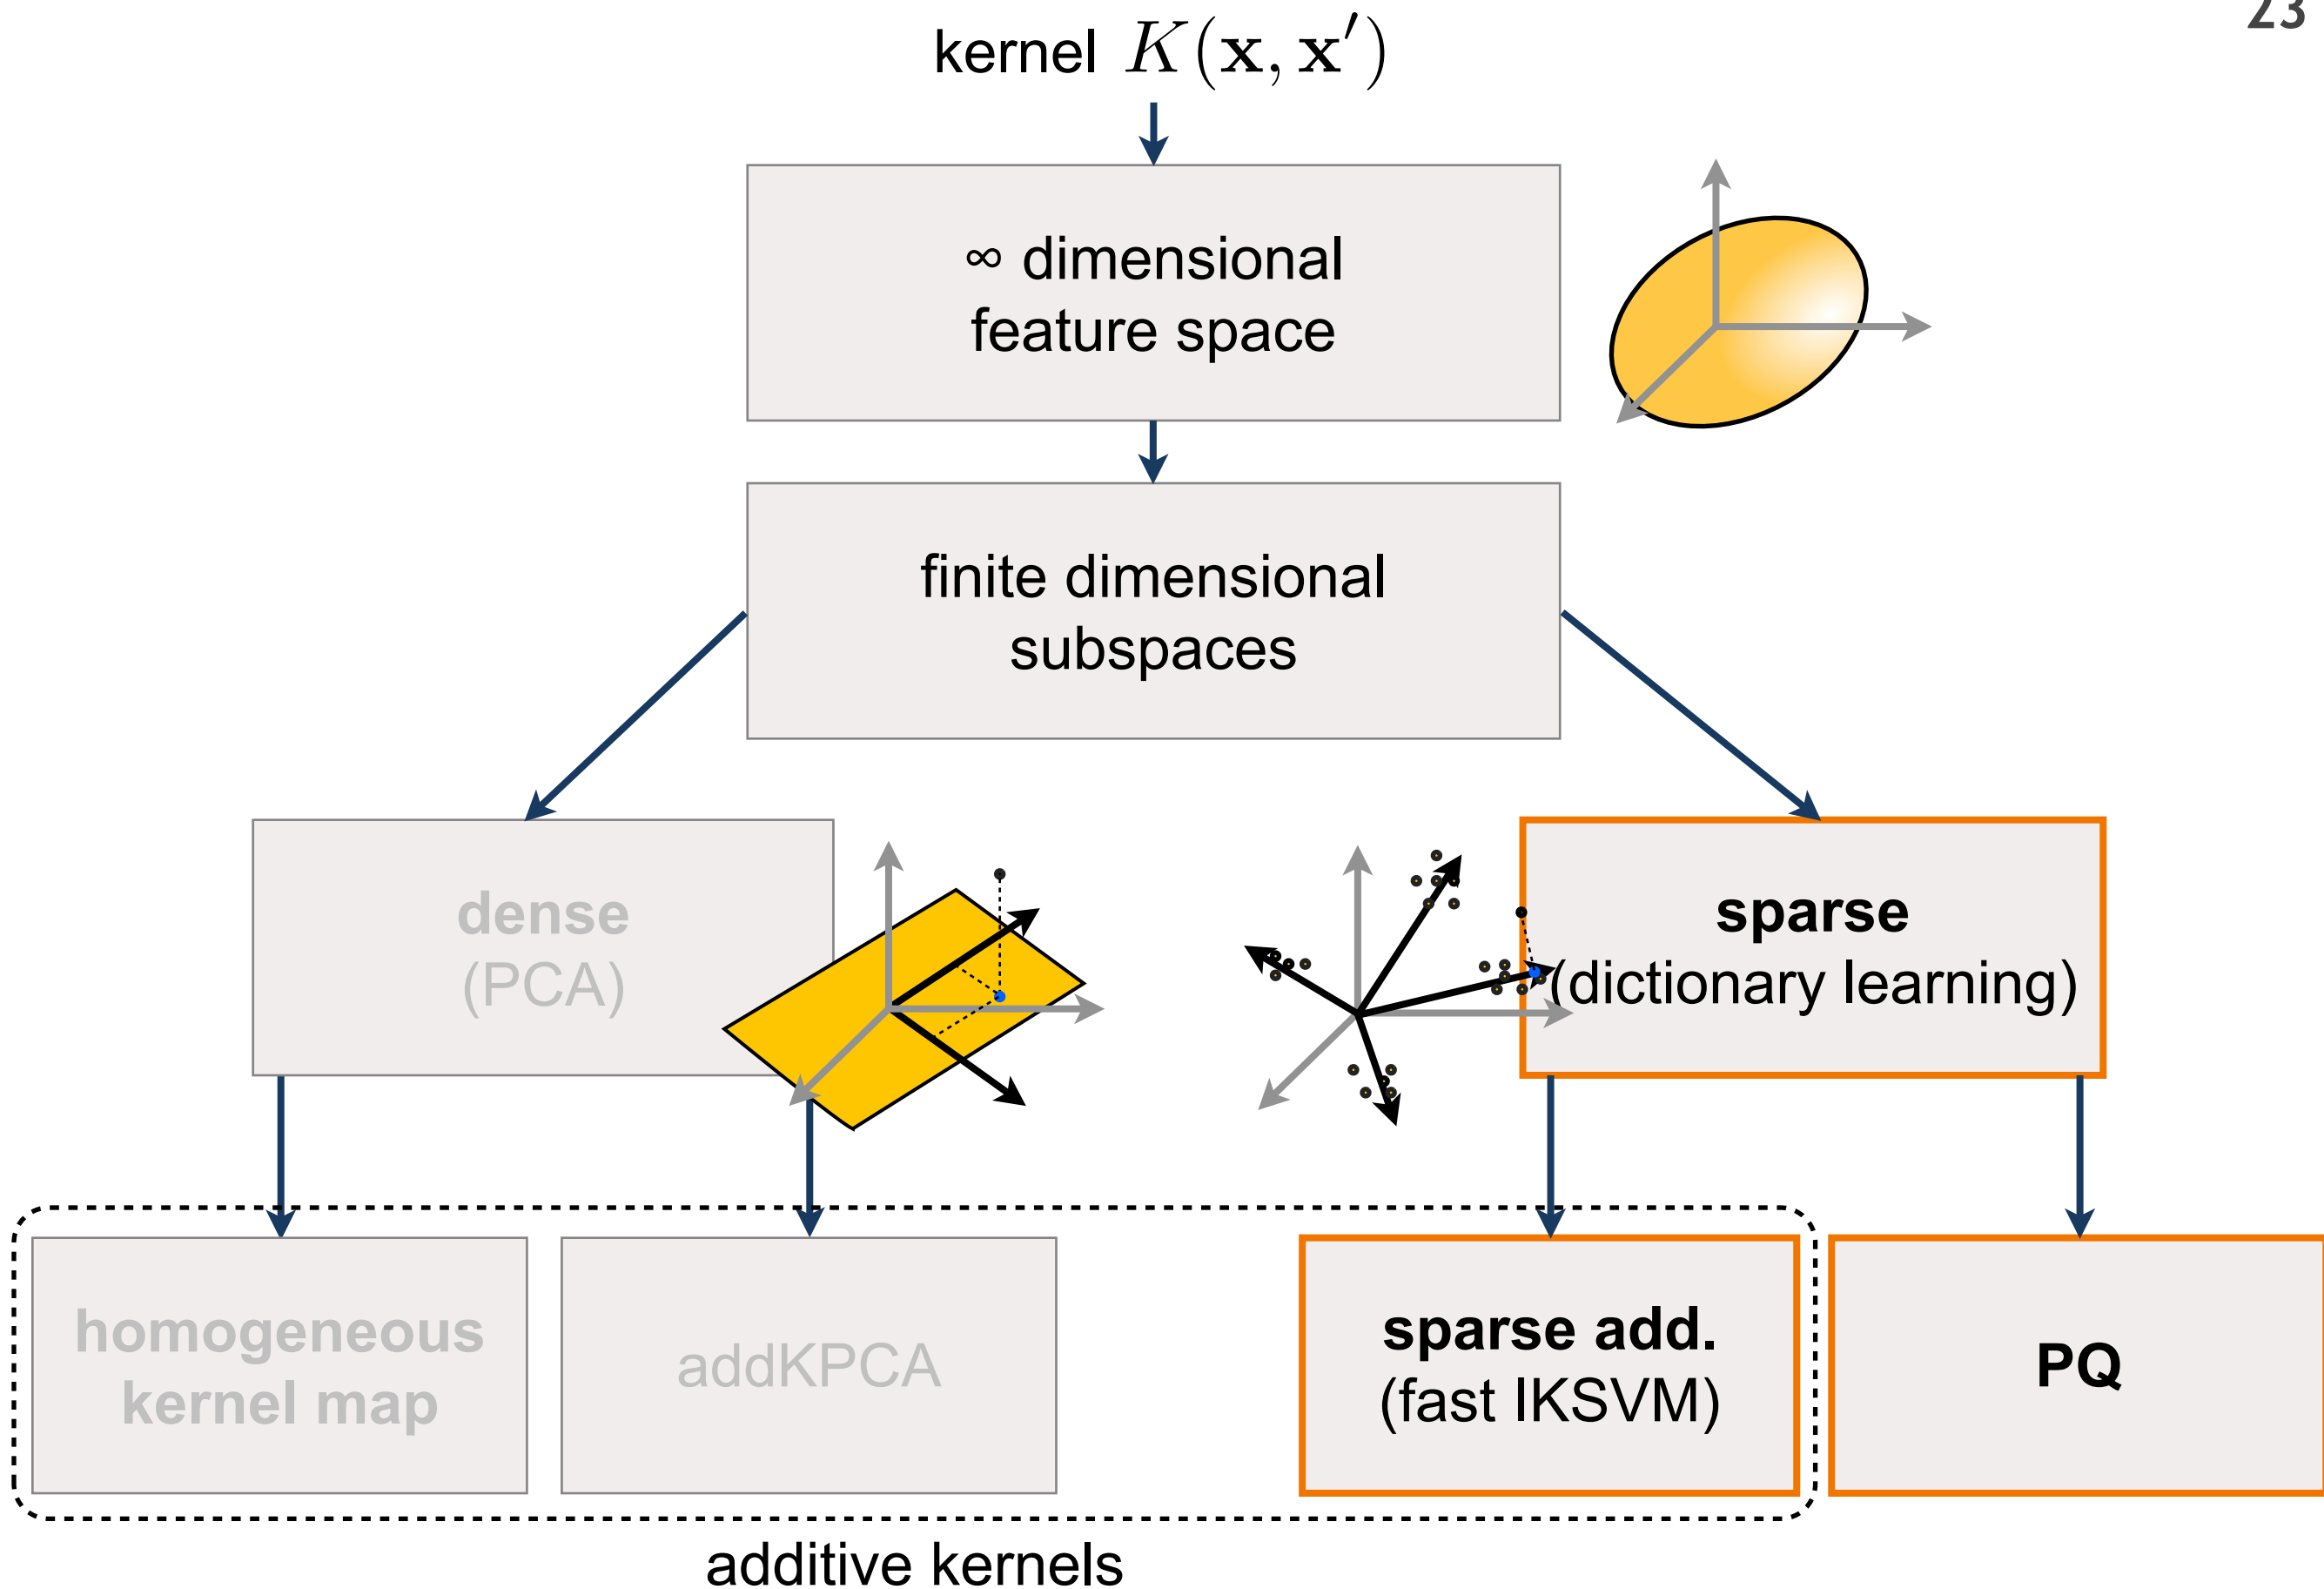


Chi2-RBF w/Random Fourier Features

```
omega = randn(10,6) ;
x = vl_homkernmap(x,1) ;
xp = vl_homkernmap(xp,1) ;
psi = [cos(omega*x) ;
        sin(omega*x)] / sqrt(10) ;
psip = [cos(omega*xp) ;
        sin(omega*xp)] / sqrt(10) ;
K = psi'*psip ;
```








Exact feature space

(e.g. reproducing kernel Hilbert space)

$$K(\mathbf{x}, \mathbf{x}') = \langle \Psi(\mathbf{x}), \Psi(\mathbf{x}') \rangle_{\mathcal{H}}$$

Data distribution

$$\{\mathbf{x}_1, \dots, \mathbf{x}_n\}$$



Sparse projection in feature space

large basis

$$Z = \{\mathbf{z}_1, \mathbf{z}_2, \dots, \mathbf{z}_D\}$$

Find best P -sparse approx. to \mathbf{x}

$$\Psi(\mathbf{x}) \approx \sum_{i=1}^D \Psi(\mathbf{z}_i) \Phi_i(\mathbf{x}) \quad \Phi(\mathbf{x}) = \begin{bmatrix} 0 \\ \alpha(\mathbf{x}) \\ 0 \\ \vdots \\ 0 \end{bmatrix}$$

✓ Arbitrarily good P -sparse approximation

● **non-diagonal** inner product



$$K(\mathbf{x}, \mathbf{x}') \approx \Phi(\mathbf{x})^\top K_{ZZ} \Phi(\mathbf{x}')$$

~~[Kernel Matching Pursuit, Vincente Bengion 02]~~

[Snelson Ghahramani 07]

[Vedaldi Zisserman 12]

$$k(x, x') \approx \Phi(x)^\top K_{ZZ} \Phi(x') = [1 \quad 0 \quad 0 \quad \dots \quad 0]$$



Fast Intersection Kernel


$$\alpha(x) = \begin{bmatrix} i + 1 - x \\ x - 1 \end{bmatrix}$$

[Maji Berg 09]

Chi2 kernel

$$\alpha(x) = \frac{2(1 + 2i)x}{(i + x)(1 + i + x)} \begin{bmatrix} i + 1 - x \\ x - 1 \end{bmatrix}$$

[Vedaldi Zisserman 12]





● \mathbf{z}_i **codewords**

$Z = \{\mathbf{z}_1, \dots, \mathbf{z}_D\}$ **codebook**


$\mathbf{x} \approx \mathbf{z}_{q(\mathbf{x})}$ **quantisation**: represent a data point by the closest codeword

saving: store only the index q , using $\log_2(D)$ bits



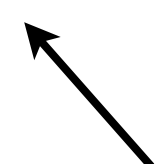


Idea
 Apply quantisation to blocks
 [Gray Neuhoff 98, Jégou et al. 11,
 Sánchez and Perronnin 11]



- **Data reconstruction** = codebook \times sparse code

$$\mathbf{x} \approx \mathbf{z}_{q(x)} = \begin{bmatrix} \mathbf{z}_1 & \dots & \mathbf{z}_D \end{bmatrix} \begin{bmatrix} 0 \\ 0 \\ \vdots \\ 1 \\ \vdots \\ 0 \end{bmatrix} = Z\Phi(\mathbf{x})$$

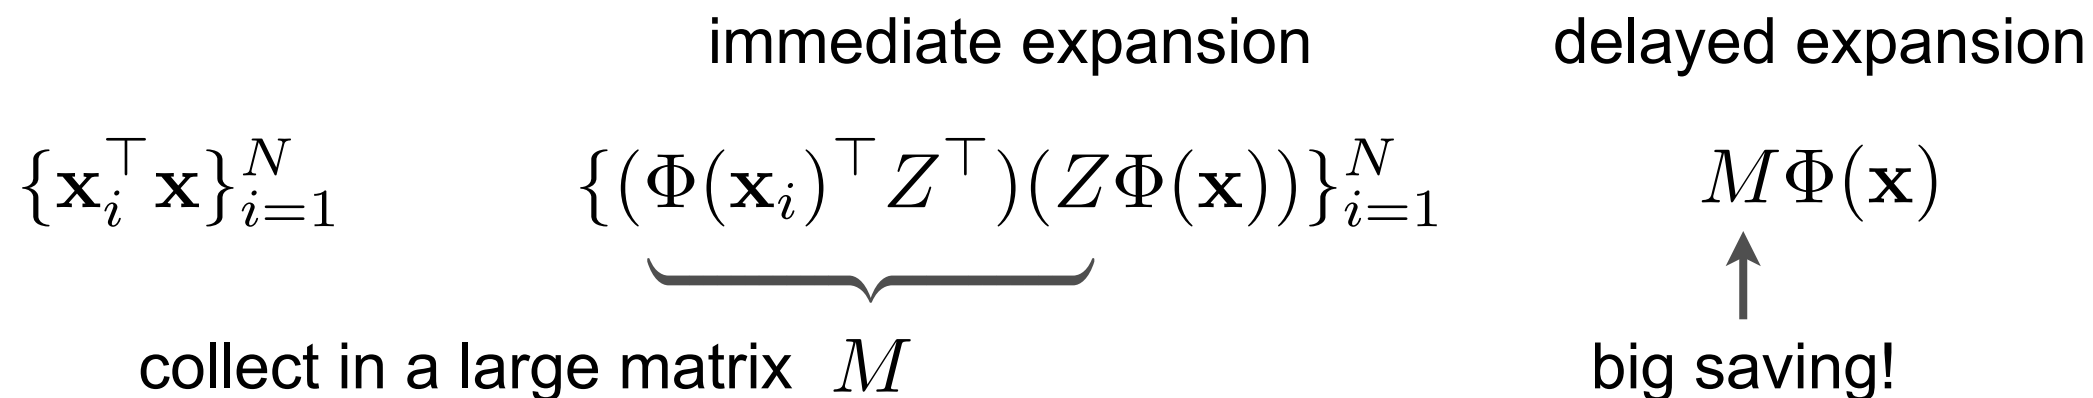
 $q(\mathbf{x})$ -th entry

- **Kernel reconstruction** = sparse feature map

$$K(\mathbf{x}, \mathbf{x}') \approx \Phi(\mathbf{x})^\top Z^\top Z \Phi(\mathbf{x}') = \Phi(\mathbf{x})^\top K_{ZZ} \Phi(\mathbf{x}')$$

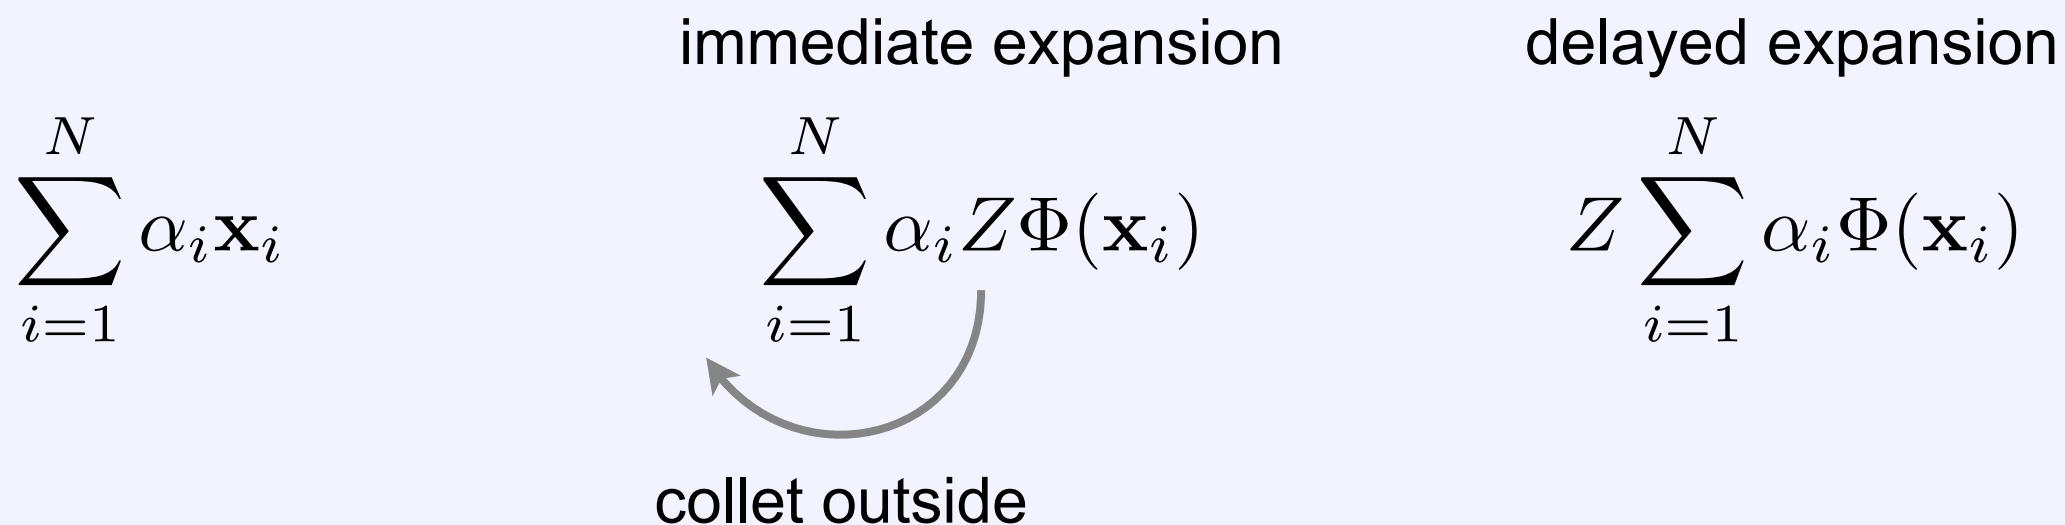
Compute quickly many **inner products**

[Jégou et al. 11]





Accumulate quickly many vectors

[Vedaldi Zisserman 12]



- 40,960-dimensional descriptors
- ~ 0.5 Mb per image
- 2GB of data

- up to 100x memory reduction
- up to 10 times faster



- **Feature maps**

- Explicit linear embeddings reproducing a kernel
- Allow tremendous speed-ups in learning
- Particularly simple & efficient for additive kernels

- **Dense low-dimensional features**

- Similar to PCA
- *Homogeneous kernel map* (analytical for homogeneous additive)
- *addKPCA* (empirical Nyström for additive)
- *Random Fourier features* (Gaussian)
- *Generalised Random Fourier Features* (Gaussian + additive)

- **Sparse high-dimensional features**

- Similar to sparse coding
- *Intersection kernel map* (sparse version)
- *Product quantisation* for compression
- Computation in the compressed domain
 - inner products and accumulation

- S. Maji and A. C. Berg. Max-margin additive classifiers for detection. In Proc. ICCV, 2009.
- S. Maji, A. C. Berg, and J. Malik. Classification using intersection kernel support vector machines is efficient. In Proc. CVPR, 2008.
- F. Perronnin, J. Sánchez, and Y. Liu. Large-scale image categorization with explicit data embedding. In Proc. CVPR, 2010.
- A. Vedaldi and A. Zisserman. Efficient additive kernels via explicit feature maps. PAMI, 34(3), 2012.
- A. Rahimi and B. Recht. Random features for large-scale kernel machines. In Proc. NIPS, 2007.
- H. Jégou, M. Douze, and C. Schmid. Product quantization for nearest neighbor search. PAMI, 33(1), 2011.
- R. M. Gray and D. L. Neuhoff. Quantization. IEEE Trans. on Information Theory, 1998.